

[/agent](#)

[/feedback](#)

[/proposal](#)

THREE DELIVERABLES

One agent, one eval, one proposal. Three sections, one site.

The agent

Luke. 27 tools, 169 turns last month from five daily-active human users (a personal-use agent at production scale, not an enterprise load). What it does, how I instrumented it, and the trace that started this.



02

`span.feedback`

Field notes from instrumenting

I started with Phoenix, switched to AX, drove both end-to-end. What pulled me over, what worked, and what broke. Headline finding: Alyx is a router, not an agent.



03

`span.proposal`

Tool Bypass: an eval primitive

A failure-class taxonomy, a precise definition, an audit of 169 production turns, four Phoenix experiments, and a continuous online task on AX scoring live spans.



Failure class:

tool bypass

Detection:

regex pre-filter
LLM-as-judge adjudicator
(running on live AX spans)

Audit sample:

169 production turns
30 days, pre-defense

Audit result:

7 real phantoms
4 false positives

Eval validation:

14 of 16 boundary cases
across Sonnet 4.6 + Opus 4.7

Live runtime:

11 of 24 AX spans scored
continuous task, zero errors

Existing Phoenix and AX evals catch hallucination and wrong-tool-selection. They don't catch an agent that asserts an action without firing the tool. I named the class *tool bypass*, wrote an eval matching Phoenix's existing module patterns, and ran it against **169 real production turns: 7 real phantoms, 4 false positives**.

The runtime story is already shipped on AX. The eval registers as a continuous online task, scores every Anthropic span automatically, and surfaces the verdict in the trace UI. The 30-day commit is about productizing the surface for AX customers, not about whether the eval works.

Bryan Plaza 2026-05-01

BSPlaza@outlook.com ↓ proposal.pdf
github.com/BSPLAZA ↗ Luke public repo
linkedin.com/in/BryanIA0 ↗ Phoenix fork