

THREE DELIVERABLES

One agent, one eval, one proposal. Three sections, one site.

01

span.agent

The agent, in production

Luke. 27 tools, 169 production turns last month, six daily-active users. What it does, how I instrumented it, and the trace that started this.

→

02

span.feedback

Field notes from instrumenting

Specific moments where Phoenix and AX worked, and where they didn't. An editorial essay rather than a feature comparison table.

→

03

span.proposal

Tool Bypass: an eval primitive

A failure-class taxonomy, a precise definition, real audit data over 169 turns, and a roadmap that names what would invalidate it.

→

Failure class:

tool bypass

Detection method:

regex first, LLM-as-judge planned

Sample:

169 production turns

30 days, post v50 defense

Result:

7 real / 4 false positives

Fork status:

private, matched to your patterns

not a PR

Existing Phoenix evals catch hallucination and wrong-tool-selection. They don't catch an agent that asserts an action without firing the tool. I named the class *tool bypass* and wrote an eval for it that matches Phoenix's existing module patterns, written as if I were already on the team.

The eval ran against **169 real production turns** from one of my agents. It caught **seven real phantom-action bugs** and surfaced **four false positives**. The false positives are useful: they argue for an LLM-as-judge follow-up rather than a regex-only approach, and they're documented row-by-row in the proposal section.

The proposal commits to an **Alyx runtime ship** in the next thirty days (with the OSS contribution as a credibility down-payment), names the dependency on handoff and memory primitives that come after, and asks the question I'd want to talk through before committing: where on the AX surface this primitive should actually live.

Bryan Plaza 2026-05-01

BSPlaza@outlook.com ↓ proposal.pdf
github.com/BSPLAZA ↗ Luke public repo
linkedin.com/in/BryanIA0 ↗ Phoenix fork